

Concurrency

Semaphore

Liu yufeng

Fx_yfliu@163.com

Hunan University

Semaphore

- **Locks** and **condition variables** can solve a broad range of relevant and interesting concurrency problems.
- **Dijkstra** and colleagues invented the **semaphore** as a single primitive for all things related to synchronization.
- One can use semaphores as both locks and condition variables.

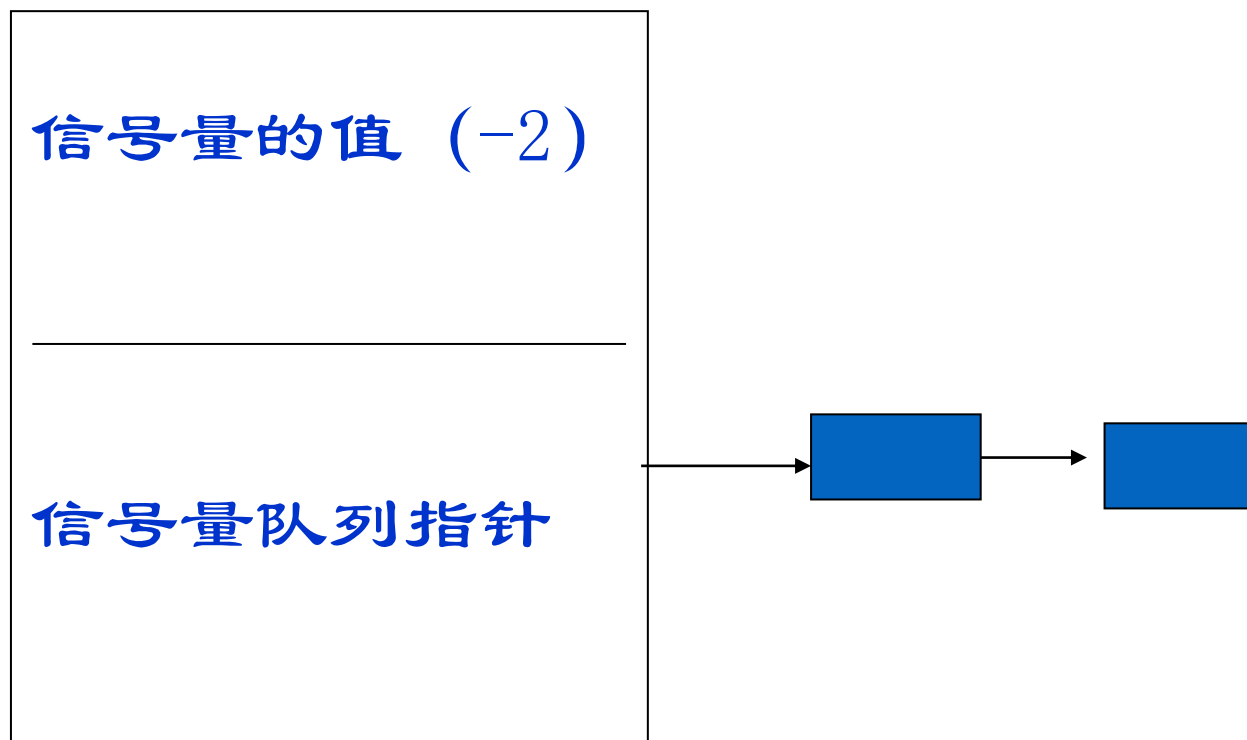
关键问题：如何使用信号量？

如何使用信号量代替锁和条件变量？什么是信号量？什么是二值信号量？用锁和条件变量来实现信号量是否简单？不用锁和条件变量，如何实现信号量？

Semaphores

- 1965年E.W.Dijkstra提出了新的同步工具--信号量和P操作（荷兰语的测试Proberen）、V操作（荷兰语的增量Verhogen）。
- 在许多操作系统的教程中，P操作又被称为wait操作，V操作有被称为signal操作

Semaphores



- Define a semaphore as a record

```
typedef struct {  
    int value;  
    struct process *Q;  
} semaphore;
```

- 信号量的数据结构:
- 两个分量: 一个是信号量的值, 另一个是信号量队列的队列指针。
- 信号量的值的含义:
- 如果 $S.value \geq 0$, 则表明系统有多少个临界资源可以分配给请求的进程。
- 如果 $S.value < 0$, 表示等待使用该临界资源的进程数。

- P和V操作原语定义:
- P(s): 将信号量s.value减1, 若结果小于0, 则调用P(s)的进程被置成等待信号量s的状态, 并把该进程挂到信号量的等待队列上。
- V(s): 将信号量s.value加1, 若结果不大于0, 则释放一个等待信号量s的进程。

对信号量整数值的修改必须是原子操作

Semaphores: A Definition

- A semaphore is as an object with **an integer value** that we can manipulate with two routines; in the POSIX standard, these routines are **`sem_wait()`** and **`sem_post()`**.

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1);
```

其中声明了一个信号量 `s`，通过第三个参数，将它的值初始化为 1。`sem_init()` 的第二个参数，在我们看到的所有例子中都设置为 0，表示信号量是在同一进程的多个线程共享的。

- `sem_wait()` 要么立刻返回（调用 `sem_wait()` 时，信号量的值大于等于 1），要么会让调用线程挂起，直到之后的一个 `post` 操作。当然，也可能多个调用线程都调用 `sem_wait()`，因此都在队列中等待被唤醒。
- `sem_post()` 增加信号量的值，如果有等待线程，唤醒其中一个。
- 当信号量的值为负数时，其绝对值就是等待线程的个数。

```
1  int sem_wait(sem_t *s) {
2      decrement the value of semaphore s by one
3      wait if value of semaphore s is negative
4  }
5
6  int sem_post(sem_t *s) {
7      increment the value of semaphore s by one
8      if there are one or more threads waiting, wake one
9  }
```

Semaphore: Definitions of Wait and Post

用P、V原语实现进程互斥

通过我们对临界区访问过程的分析，信号量机制中P原语相当于进入区操作，V原语相当于退出区操作。

用P、V原语实现进程互斥就是将**临界区**置于**P**和**V**两个原语操作之间。

进入时执行P操作，使信号量S.value减1，如果 $S.value \geq 0$ ，则进入临界区，否则不可进入。进程退出临界区时，执行V操作，使信号量 $S.value + 1$ ，表示释放临界资源。



Binary Semaphores (Locks)

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize semaphore to X; what should X be?  
3  
4  sem_wait(&m);  
5  // critical section here  
6  sem_post(&m);
```

- Critical to making this work, though, is the **initial value of the semaphore *m*** (initialized to **X** in the figure).
What should **X** be?

表 31.1

追踪线程：单线程使用一个信号量

信号量的值	线程 0	线程 1
1		
1	调用 sem_wait()	
0	sem_wait() 返回	
0	(临界区)	
0	调用 sem_post()	
1	sem_post() 返回	

表 31.2

追踪线程：两个线程使用一个信号量

值	线程 0	状态	线程 1	状态
1		运行		就绪
1	调用 sem_wait()	运行		就绪
0	sem_wait()返回	运行		就绪
0	(临界区：开始)	运行		就绪
0	中断；切换到→T1	就绪		运行
0		就绪	调用 sem_wait()	运行
-1		就绪	sem 减 1	运行
-1		就绪	(sem<0) →睡眠	睡眠
-1		运行	切换到→T0	睡眠
-1	(临界区：结束)	运行		睡眠
-1	调用 sem_post()	运行		睡眠
0	增加 sem	运行		睡眠
0	唤醒 T1	运行		就绪
0	sem_post()返回	运行		就绪
0	中断；切换到→T1	就绪		运行
0		就绪	sem_wait()返回	运行
0		就绪	(临界区)	运行
0		就绪	调用 sem_post()	运行
1		就绪	sem_post()返回	运行

- Note that **Thread 1** goes into the sleeping state when it tries to acquire the already-held lock; only when **Thread 0** runs again can **Thread 1** be awoken and potentially run again.

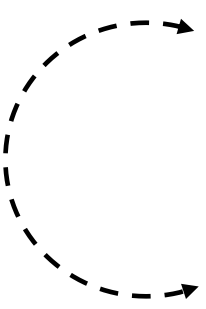
Critical Section of n Processes

- Shared data:

semaphore mutex; //initially mutex = 1

- Process P_i :

```
do {  
    P(mutex);  
    critical section  
    V(mutex);  
    remainder section  
} while (1);
```




锁，互斥

Semaphore as a General Synchronization Tool

- Execute B in P_j only after A executed in P_i

- Code:



Semaphores As Condition Variables

- Semaphores are also useful when a thread wants to halt its progress waiting for a condition to become true.

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     Pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

A Parent Waiting For Its Child

- What should the **initial value** of this semaphore be?

表 31.3

追踪线程：父线程等待子线程（场景 1）

值	父线程	状态	子线程	状态
0	create(子线程)	运行	(子线程产生)	就绪
0	调用 sem_wait()	运行		就绪
-1	sem 减 1	运行		就绪
-1	(sem<0)→ 睡眠	睡眠		就绪
-1	切换到→子线程	睡眠	子线程运行	运行
-1		睡眠	调用 sem_post()	运行
0		睡眠	sem 增 1	运行
0		就绪	wake(父线程)	运行
0		就绪	sem_post()返回	运行
0		就绪	中断：切换到→父线程	就绪
0	sem_wait()返回	运行		就绪

Thread Trace: Parent Waiting For Child (Case 1)

```

1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     Pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```


表 31.4

追踪线程：父线程等待子线程（场景 2）

值	父线程	状态	子线程	状态
0	create（子线程）	运行	(子线程产生)	就绪
0	中断：切换到→子线程	就绪	子线程运行	运行
0		就绪	调用 sem_post()	运行
1		睡眠	sem 增 1	运行
1		就绪	wake(没有线程)	运行
1		就绪	sem_post()返回	运行
1	父线程运行	运行	中断：切换到→父线程	就绪
1	调用 sem_wait()	运行		就绪
0	sem 减 1	运行		就绪
0	(sem>=0)→不用睡眠	运行		就绪
0	sem_wait()返回	运行		就绪

```


1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     Pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }

```

生产者-消费者问题表述

把并发进程的互斥和同步问题一般化，我们可以得出一个抽象化的一般模型，即生产者-消费者问题。

- 有 n 个生产者和 m 个消费者，连接在一个有 k 个单位缓冲区的有界缓冲上。其中， p_i 和 c_j 都是并发进程，只要缓冲区未满，生产者 p_i 生产的产品就可投入缓冲区；只要缓冲区不空，消费者进程 c_j 就可从缓冲区取走并消耗产品。



The Producer/Consumer (Bounded-Buffer) Problem

- Our first attempt at solving the problem introduces **two semaphores**, **empty** and **full**, which the threads will use to indicate when a buffer entry has been emptied or filled, respectively.

```

1  int buffer[MAX];
2  int fill = 0;
3  int use  = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;    // line g2
13     return tmp;
14 }
15
16 sem_t empty;
17 sem_t full;
18
19 void *producer(void *arg) {
20     int i;
21     for (i = 0; i < loops; i++) {
22         sem_wait(&empty);    // line P1
23         put(i);              // line P2
24         sem_post(&full);     // line P3
25     }
26 }
27
28 void *consumer(void *arg) {
29     int i, tmp = 0;
30     while (tmp != -1) {
31         sem_wait(&full);    // line C1
32         tmp = get();        // line C2
33         sem_post(&empty);   // line C3
34         printf("%d\n", tmp);
35     }
36 }
37
38 int main(int argc, char *argv[]) {
39     // ...
40     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
41     sem_init(&full, 0, 0);    // ... and 0 are full
42     // ...
43 }

```

Let us first imagine that **MAX=1** (there is only one buffer in the array), and see if this works.

Imagine again there are **two threads**, a producer and a consumer.

```

1  int buffer[MAX];
2  int fill = 0;
3  int use  = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;    // line g2
13     return tmp;
14 }
15
16 sem_t empty;
17 sem_t full;
18
19 void *producer(void *arg) {
20     int i;
21     for (i = 0; i < loops; i++) {
22         sem_wait(&empty);      // line P1
23         put(i);                // line P2
24         sem_post(&full);       // line P3
25     }
26 }
27
28 void *consumer(void *arg) {
29     int i, tmp = 0;
30     while (tmp != -1) {
31         sem_wait(&full);       // line C1
32         tmp = get();           // line C2
33         sem_post(&empty);      // line C3
34         printf("%d\n", tmp);
35     }
36 }
37
38 int main(int argc, char *argv[]) {
39     // ...
40     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
41     sem_init(&full, 0, 0);    // ... and 0 are full
42     // ...
43 }

```

假设 **consumer** 先运行. 阻塞在C1 (full从0 变为-1), 等待另一个线程调用 sem_post(&full).

假设**producer**随后执行. 执行到P1 (empty 为 1), 继续执行到 P2, 向缓冲区放入数据, 然后执行P3调用 sem_post(&full), 唤醒consumer.

这时可能有两种情况. 第一, producer继续执行多次循环到P1并阻塞(empty为0); 第二, consumer开始执行调用sem_wait(&full)(C1),取数据. 这两种情况都**正常** (包括多个 **consumer**和**producer**).

```

1  int buffer[MAX];
2  int fill = 0;
3  int use  = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;    // line g2
13     return tmp;
14 }

1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);    // line P1
8          put(i);              // line P2
9          sem_post(&full);     // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);      // line C1
17         tmp = get();          // line C2
18         sem_post(&empty);     // line C3
19         printf("%d\n", tmp);
20     }
21 }
22
23 int main(int argc, char *argv[]) {
24     // ...
25     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
26     sem_init(&full, 0, 0);    // ... and 0 are full
27     // ...
28 }

```

假设**MAX**大于1, 并且假设
 有多个producer和
 consumer. 问题来了: **a**
race condition. 哪里?

假设两个producer(P_a 和 P_b)同
 时调用put(). 假设 P_a 先执
 行f1, 但在执行f2之前被中
 断, P_b 开始执行, 也在f1处写
 入数据, 这将覆盖 P_a 写入的
 数据.

A Solution: Adding Mutual Exclusion

- What we've forgotten here is **mutual exclusion**.
- The filling of a buffer and incrementing of the index into the buffer is a **critical section**, and thus must be guarded carefully.

```

1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);          // line p0 (NEW LINE)
9          sem_wait(&empty);          // line p1
10         put(i);                    // line p2
11         sem_post(&full);            // line p3
12         sem_post(&mutex);          // line p4 (NEW LINE)
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&mutex);          // line c0 (NEW LINE)
20         sem_wait(&full);           // line c1
21         int tmp = get();           // line c2
22         sem_post(&empty);          // line c3
23         sem_post(&mutex);          // line c4 (NEW LINE)
24         printf("%d\n", tmp);
25     }
26 }
27
28 int main(int argc, char *argv[]) {
29     // ...
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
31     sem_init(&full, 0, 0);    // ... and 0 are full
32     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock (NEW LINE)
33     // ...
34 }

```

Adding Mutual Exclusion (Incorrectly)

- Add some **locks around the entire put()/get()** parts of the code, as indicated by the **NEW LINE** comments. Seems like the right idea, but it also **doesn't work**. **Why?**
- Why does deadlock occur? What sequence of steps must happen for the program to deadlock?

Avoiding Deadlock

- 假设有两个线程，一个生产者和一个消费者。消费者先运行获得锁(c0)，然后对 full 信号量执行 sem_wait() (c1)。因为还没有数据，所以消费者阻塞。但是，重要的是，此时消费者仍然持有锁。
- 然后生产者运行，它首先对二值互斥信号量调用 sem_wait()(p0)。锁已经被持有，因此生产者也被卡住。
- 这里出现了一个循环等待。消费者持有互斥量，等待在 full 信号量上。生产者可以发送 full 信号，却在等待互斥量。因此，生产者和消费者互相等待对方——典型的死锁。

```


1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&empty);          // line p1
9          sem_wait(&mutex);          // line p1.5 (MOVED MUTEX HERE...)
10         put(i);                    // line p2
11         sem_post(&mutex);          // line p2.5 (... AND HERE)
12         sem_post(&full);           // line p3
13     }
14 }
15
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&full);            // line c1
20         sem_wait(&mutex);          // line c1.5 (MOVED MUTEX HERE...)
21         int tmp = get();            // line c2
22         sem_post(&mutex);          // line c2.5 (... AND HERE)
23         sem_post(&empty);          // line c3
24         printf("%d\n", tmp);
25     }
26 }
27
28 int main(int argc, char *argv[]) {
29     // ...
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with...
31     sem_init(&full, 0, 0);    // ... and 0 are full
32     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock
33     // ...
34 }

```

Adding Mutual Exclusion (Correctly)

- To solve this problem, we must **reduce the scope of the lock**. We simply move the mutex to be just around the critical section; the full and empty wait and signal code is left outside.

多个生产者、多个消费者、共享多个缓冲区的解



- a. 设置公有信号量**mutex**，以实现互斥；
- b. 设置私有信号量**empty**和**full**，以实现同步；
- c. 赋初值：**empty=k, full=0, mutex=1**；
- d. 实现P、V操作。

Bounded-Buffer Problem

Shared data

semaphore full, empty, mutex;

Initially: full = 0, empty = n, mutex = 1

Producer

do {

...

produce an item in nextp

...

wait(empty);

wait(mutex);

...

add nextp to buffer

...

signal(mutex);

signal(full);

} while (1);

Consumer

do {

wait(full)

wait(mutex);

...

remove an item from buffer to nextc

...

signal(mutex);

signal(empty);

...

consume the item in nextc

...

} while (1);

wait (empty) 操作和
wait (mutex) 操作的
顺序互换, 以及
wait (full) 操作和
wait (mutex) 操作的
顺序互换, 可能会怎样?

为什么需要wait (mutex)

假定：此时 $\text{mutex}=1$ ， $\text{full}=\text{n}$ ， $\text{empty}=0$ ，并且生产者先占用CPU



死锁！

交换`signal`的顺序不会有任何影响

Producer

```
do {  
    ...  
    produce an item in nextp  
    ...  
    wait(mutex);  
    wait(empty);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```

Consumer

```
do {  
    wait(mutex)  
    wait(full);  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```

还有一种死锁的可能性是什么？

P、V操作小结

1) 信号量的物理含义:

$S > 0$: 表示有 S 个资源可用

$S = 0$: 表示无资源可用

$S < 0$: 则 $|S|$ 表示 S 等待队列中的进程个数

$P(S)$: 表示申请一个资源

$V(S)$: 表示释放一个资源。

P、V操作小结

2) P、V操作必须成对出现，有一个P操作就一定有一个V操作

当为互斥操作时，它们同处于同一进程

当为同步操作时，则不在同一进程中出现

如果P(S1)和P(S2)两个操作在一起，那么P操作的顺序至关重要，一个同步P操作与一个互斥P操作在一起时，同步P操作在互斥P操作前

而两个V操作无关紧要

P、V操作小结

3) P.V操作的优缺点

优点:


简单, 而且表达能力强 (用P.V操作可解决任何同步互斥问题)

缺点:

不够安全; P.V操作使用不当会出现死锁;
遇到复杂同步互斥问题时实现复杂

读者-写者问题

- 多个并发进程共享一个数据对象。其中有些进程只是读取这些共享数据的内容，而其它进程可能想要更新共享对象。我们称只是要读取共享对象的进程为读者，对共享对象进行更新操作的进程为写者，以此来区别对待这两种进程类型。



读者-写者问题

制约条件分析:

- 1、允许多个进程同时读文件（读-读允许）；
- 2、不允许在进程读文件时让另外一进程去写文件；
有进程在写文件时不让另外一个进程去读该文件
（“读-写”互斥）；
- 3、不允许多个写进程同时写同一文件（“写-写”互斥）。

读者-写者问题

因为允许多个进程同时读，系统应记录读进程的个数，而每个读进程去读文件或读文件结束后都要修改读者个数，因此，读者个数又是若干读进程的共享变量，即软件临界资源，它们也必须互斥地修改这个变量。

综上，我们定义：

- 1、**wrt**：写互斥信号量，用于“读-写”和“写-写”互斥；初值为1；
- 2、公共变量**readcount**用于记录当前正在读文件的读者数目，初值为0；
- 3、**mutex**：用于若干读进程对读者个数修改的互斥，初值为1。

Shared data:

semaphore mutex, wrt;

Initially

mutex = 1, wrt = 1, readcount = 0

Reader:

wait(mutex);

readcount++;

if (readcount == 1)

wait(wrt);

signal(mutex);

...

reading is performed

...

wait(mutex);

readcount--;

if (readcount == 0)

signal(wrt);

signal(mutex);

Writer:

wait(wrt);

...

writing is performed

...

signal(wrt);

满足写写互斥吗?

满足读读允许吗?

满足读写互斥吗?

假定先有写进程，然后读进程想要读。

假定先有读进程，然后写进程想要写。

现在是完美的解决方案了吗?

- 读者优先
- 只要不断的有读者来读，那么readcount就一直会大于0，那么永远不会触发signal(wrt)条件，从而导致写者饥饿。
- 该问题被称为**第一读者优先问题**

写者优先：（作业）

- 1.写者线程的优先级高于读者线程。
- 2.当写者到来时，只有那些已经获得授权的读进程才被允许完成它们的操作，写者之后到来的读者将被推迟，直到写者完成。
- 3.当没有写者进程时读者进程应该能够同时读取文件。

公平竞争：（作业）

- 1.优先级相同。
- 2.写者、读者互斥访问。
- 3.只能有一个写者访问临界区。
- 4.可以有多个读者同时访问临界资源。


Reader-Writer Locks

```
1  typedef struct _rwlock_t {
2      sem_t lock;          // binary semaphore (basic lock)
3      sem_t writelock;     // allow ONE writer/MANY readers
4      int    readers;      // #readers in critical section
5  } rwlock_t;
6
7  void rwlock_init(rwlock_t *rw) {
8      rw->readers = 0;
9      sem_init(&rw->lock, 0, 1);
10     sem_init(&rw->writelock, 0, 1);
11 }
12
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     rw->readers++;
16     if (rw->readers == 1) // first reader gets writelock
17         sem_wait(&rw->writelock);
18     sem_post(&rw->lock);
19 }
20
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0) // last reader lets it go
25         sem_post(&rw->writelock);
26     sem_post(&rw->lock);
27 }
28
29 void rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32
33 void rwlock_release_writelock(rwlock_t *rw) {
34     sem_post(&rw->writelock);
35 }
```

读者很容易
饿死写者。

The Dining Philosophers

- One of the most famous concurrency problems posed, and solved, by **Dijkstra**.
- The problem is famous because it is fun and somewhat **intellectually interesting**; however, **its practical utility is low**.



The Dining Philosophers

- Assume there are five “philosophers” sitting around a table.
- Between each pair of philosophers is a single fork (**five total**).
- The philosophers each have times where they **think**, and don’t need any forks, and times where they **eat**.
- In order to eat, a philosopher **needs two forks**, **both** the one on their **left** and the one on their **right**.

- Here is the basic loop of each philosopher:

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

关键的挑战就是如何实现 `getforks()` 和 `putforks()` 函数，保证没有死锁，没有哲学家饿死，并且并发度更高

- We' ll use a few helper functions to get us towards a solution:

```
int left(int p)    { return p; }  
int right(int p)  { return (p + 1) % 5; }
```

- 我们需要一些信号量来解决这个问题。假设需要 5 个，每个餐叉一个: `sem_t forks[5]`.

Broken Solution

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

```
int left(int p) { return p; }  
int right(int p) { return (p + 1) % 5; }
```

```
sem_t forks[5]
```

```
1 void getforks() {  
2     sem_wait(forks[left(p)]);  
3     sem_wait(forks[right(p)]);  
4 }  
5  
6 void putforks() {  
7     sem_post(forks[left(p)]);  
8     sem_post(forks[right(p)]);  
9 }
```

- 我们把每个信号量（在 fork 数组中）都用 1 初始化。同时假设每个哲学家知道自己的编号（p）。我们可以写出 getforks() 和 putforks() 函数。

- 为了拿到餐叉，我们依次获取每把餐叉的锁——先是左手边的，然后是右手边的。结束就餐时，释放掉锁

- Simple, no? Unfortunately, in this case, simple means broken. **Can you see the problem that arises?**

- The problem is **deadlock**. If each philosopher happens to **grab the fork on their left** before any philosopher can grab the fork on their right, each will be **stuck holding one fork** and **waiting for another, forever**.

死锁预防策略

死锁解决方法：

至多只允许四位哲学家同时去拿左边的筷子；

规定奇数号哲学家先拿起他左边的筷子，而偶数号哲学家先拿起他右边的筷子。

仅当哲学家左右两边的筷子均可用时才允许他拿起筷子；

A Solution: Breaking The Dependency

- The simplest way to attack this problem is to **change how forks are acquired** by at least one of the philosophers; indeed, **this is how Dijkstra himself solved the problem.**

```
1  void getforks() {
2      if (p == 4) {
3          sem_wait(forks[right(p)]);
4          sem_wait(forks[left(p)]);
5      } else {
6          sem_wait(forks[left(p)]);
7          sem_wait(forks[right(p)]);
8      }
9  }
```

- Because the last philosopher tries to grab right before left, there is no situation where each philosopher grabs one fork and is stuck waiting for another; **the cycle of waiting is broken.**

How To Implement Semaphores

```
1  typedef struct __Zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  } Zem_t;
6
7  // only one thread can call this
8  void Zem_init(Zem_t *s, int value) {
9      s->value = value;
10     Cond_init(&s->cond);
11     Mutex_init(&s->lock);
12 }
13
14 void Zem_wait(Zem_t *s) {
15     Mutex_lock(&s->lock);
16     while (s->value <= 0)
17         Cond_wait(&s->cond, &s->lock);
18     s->value--;
19     Mutex_unlock(&s->lock);
20 }
21
22 void Zem_post(Zem_t *s) {
23     Mutex_lock(&s->lock);
24     s->value++;
25     Cond_signal(&s->cond);
26     Mutex_unlock(&s->lock);
27 }
```

实现的 Zemaphore 和 Dijkstra 定义的信号量有一点细微区别，就是我们没有保持当信号量的值为负数时，让它反映出等待的线程数

- **Curiously**, building locks and condition variables out of semaphores is **a much trickier proposition**. Some highly experienced concurrent programmers tried to do this in the Windows environment, and **many different bugs ensued** [B04].

[B04] “Implementing Condition Variables with Semaphores”

Andrew Birrell

December 2004

An interesting read on how difficult implementing CVs on top of semaphores really is, and the mistakes the author and co-workers made along the way. Particularly relevant because the group had done a ton of concurrent programming; Birrell, for example, is known for (among other things) writing various thread-programming guides.

End